

# Tennis

June 29, 2020

## 1 Collaboration and Competition

---

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

### 1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.0
```

The environment is already saved in the Workspace and can be accessed at the file path provided below.

```
In [2]: from unityagents import UnityEnvironment
import numpy as np
```

```
env = UnityEnvironment(file_name="/data/Tennis_Linux_NoVis/Tennis")
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
    Number of Brains: 1
```

```
    Number of External Brains : 1
```

```
    Lesson number : 0
```

```
    Reset Parameters :
```

```
Unity brain name: TennisBrain
```

```
    Number of Visual Observations (per agent): 0
```

```
    Vector Observation space type: continuous
```

```
    Vector Observation space size (per agent): 8
```

```
    Number of stacked Vector Observation: 3
```

```
Vector Action space type: continuous
Vector Action space size (per agent): 2
Vector Action descriptions: ,
```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```
In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]
```

### 1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```
In [4]: # reset the environment
env_info = env.reset(train_mode=True)[brain_name]

# number of agents
num_agents = len(env_info.agents)
print('Number of agents:', num_agents)

# size of each action
action_size = brain.vector_action_space_size
print('Size of each action:', action_size)

# examine the state space
states = env_info.vector_observations
state_size = states.shape[1]
print('There are {} agents. Each observes a state with length: {}'.format(states.shape[0], state_size))
print('The state for the first agent looks like:', states[0])
```

```

Number of agents: 2
Size of each action: 2
There are 2 agents. Each observes a state with length: 24
The state for the first agent looks like: [ 0.          0.          0.          0.          0.
 0.          0.          0.          0.          0.          0.          0.
 0.          0.          -6.65278625 -1.5          -0.          0.
 6.83172083  6.          -0.          0.          ]

```

### 1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agents while they are training**, and you should set `train_mode=True` to restart the environment.

```

In [6]: for i in range(5):
        env_info = env.reset(train_mode=False)[brain_name]
        states = env_info.vector_observations
        scores = np.zeros(num_agents)
        counter = 0
        while True:
            counter += 1
            actions = np.random.randn(num_agents, action_size)
            actions = np.clip(actions, -1, 1)
            env_info = env.step(actions)[brain_name]
            next_states = env_info.vector_observations
            rewards = env_info.rewards
            dones = env_info.local_done

            scores += env_info.rewards
            states = next_states
            if np.any(dones):
                break
        print('Total score (averaged over agents) this episode: {} {}'.format(np.mean(scores), counter))

Total score (averaged over agents) this episode: -0.004999999888241291 15
Total score (averaged over agents) this episode: -0.004999999888241291 16
Total score (averaged over agents) this episode: -0.004999999888241291 14
Total score (averaged over agents) this episode: -0.004999999888241291 13
Total score (averaged over agents) this episode: -0.004999999888241291 21

```

#### 1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agents while they are training. However, *after training the agents*, you can download the saved model weights to watch the agents on your own machine!

```

In [6]: from maddpg import MADDPGAgent
        from collections import deque
        import matplotlib.pyplot as plt
        import torch

        o_dim = 24
        a_dim = 2

```

```

x_dim = num_agents * o_dim

LR_ACTOR = 1e-2
LR_CRITIC = 1e-4
LEARN_NUM = 1
BATCH_SIZE = 128
GAMMA = .99
TAU = 0.001
BUFFER_SIZE = int(1e6)

eps_start = 1.0
eps_end = 0.01
eps_decay = 0.99

maddpg = MADDPGAgent(num_agents, x_dim, o_dim, a_dim,
                      lr_actor = LR_ACTOR, lr_critic = LR_CRITIC, learn_num = LEARN_NUM,
                      batch_size = BATCH_SIZE, gamma = GAMMA, tau = TAU,
                      buffer_size = BUFFER_SIZE, seed = 1234)

NUM_EPISODES = 2500
PRINT_EVERY = 100

In [7]: # train agents

score_list = []
score_window = deque(maxlen = PRINT_EVERY)

eps = eps_start
best_score = 0.5
for e in range(1, NUM_EPISODES+1):
    env_info = env.reset(train_mode=True)[brain_name]    # reset environment
    states = env_info.vector_observations
    scores = np.zeros(num_agents)
    while True:
        actions = maddpg.get_actions(states, eps)

        env_info = env.step(actions)[brain_name]
        next_states = env_info.vector_observations
        dones = env_info.local_done
        rewards = env_info.rewards

        maddpg.step(states, actions, rewards, next_states, dones )

        states = next_states
        scores += rewards
        if np.any(dones):
            break

```

```

eps = max(eps_end, eps_decay*eps) # decrease exploartion
max_score = np.max(scores) # max score among the agents

score_window.append(max_score)
mean_score = np.mean(score_window)
score_list.append([max_score, mean_score])

maddpg.lr_scheduler()

print('\rEpisode {} \t score {:.2f} \t mean score {:.2f}'.format(e, max_score, mean_
if e % PRINT_EVERY == 0:
    print('\rEpisode {} \t score {:.2f} \t mean score {:.2f}'.format(e, max_score, m


if mean_score >= best_score:
    print('\rEpisode {} \t score {:.2f} \t mean score {:.2f} \t saved!'.format(e, max_
    best_score = mean_score
    for i, agent in enumerate(maddpg.agents):
        torch.save(agent.actor.state_dict(), 'checkpoint_actor'+str(i)+'.pth')
        torch.save(agent.critic.state_dict(), 'checkpoint_critic'+str(i)+'.pth')

```

Episode 100	score	0.00	mean score	0.00	
Episode 200	score	0.00	mean score	0.01	
Episode 300	score	0.10	mean score	0.03	
Episode 400	score	0.00	mean score	0.01	
Episode 500	score	0.00	mean score	0.00	
Episode 600	score	0.00	mean score	0.00	
Episode 700	score	0.00	mean score	0.00	
Episode 800	score	0.00	mean score	0.01	
Episode 900	score	0.00	mean score	0.02	
Episode 1000	score	0.00	mean score	0.02	
Episode 1100	score	0.09	mean score	0.02	
Episode 1200	score	0.10	mean score	0.04	
Episode 1300	score	0.10	mean score	0.05	
Episode 1400	score	0.09	mean score	0.07	
Episode 1500	score	0.40	mean score	0.13	
Episode 1600	score	0.70	mean score	0.15	
Episode 1700	score	0.80	mean score	0.14	
Episode 1800	score	0.20	mean score	0.32	
Episode 1900	score	0.10	mean score	0.26	
Episode 2000	score	0.10	mean score	0.49	
Episode 2010	score	0.70	mean score	0.50	saved!
Episode 2012	score	1.00	mean score	0.51	saved!
Episode 2016	score	0.90	mean score	0.52	saved!
Episode 2018	score	0.60	mean score	0.52	saved!
Episode 2019	score	0.20	mean score	0.52	saved!
Episode 2020	score	0.10	mean score	0.52	saved!
Episode 2021	score	0.20	mean score	0.52	saved!

Episode 2022	score	0.30	mean score	0.52	saved!
Episode 2023	score	0.20	mean score	0.53	saved!
Episode 2024	score	0.50	mean score	0.53	saved!
Episode 2025	score	0.50	mean score	0.53	saved!
Episode 2026	score	0.50	mean score	0.53	saved!
Episode 2027	score	0.50	mean score	0.54	saved!
Episode 2028	score	1.40	mean score	0.54	saved!
Episode 2029	score	0.60	mean score	0.55	saved!
Episode 2030	score	0.40	mean score	0.55	saved!
Episode 2032	score	1.20	mean score	0.56	saved!
Episode 2034	score	1.70	mean score	0.57	saved!
Episode 2035	score	0.20	mean score	0.57	saved!
Episode 2100	score	0.20	mean score	0.30	
Episode 2200	score	0.10	mean score	0.20	
Episode 2300	score	0.20	mean score	0.20	
Episode 2400	score	0.10	mean score	0.16	
Episode 2500	score	0.10	mean score	0.17	

```
In [8]: fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(score_list)), score_list)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()
```



```

In [10]: # watch trained agents
         for i, agent in enumerate(maddpg.agents):
             agent.actor.load_state_dict(torch.load('checkpoint_actor'+str(i)+'.pth'))

         for i in range(5):
             env_info = env.reset(train_mode=False)[brain_name]
             states = env_info.vector_observations
             scores = np.zeros(num_agents)
             while True:
                 actions = maddpg.get_actions(states, eps)
                 env_info = env.step(actions)[brain_name]
                 next_states = env_info.vector_observations
                 rewards = env_info.rewards
                 dones = env_info.local_done

                 scores += env_info.rewards
                 states = next_states
                 if np.any(dones):
                     break
             print('Total score (averaged over agents) this episode: {}'.format(np.mean(scores)))

Total score (averaged over agents) this episode: 0.7950000120326877
Total score (averaged over agents) this episode: 0.04500000085681677
Total score (averaged over agents) this episode: 0.44500000681728125
Total score (averaged over agents) this episode: 0.04000000096857548
Total score (averaged over agents) this episode: 0.1450000023469329

```

When finished, you can close the environment.

```

In [11]: env.close()

```

```

In [ ]:

```